

CSE 4020 Machine Learning

Digital Assignment

Sujay Kumar M 20BDS0294

Computer Science Engineering with Specialization with DataScience

sujaykumarreddy.m2020@vitstudent.ac.in

https://github.com/sujaykumarmag/CSE4020_digital_assignment

April 9, 2023

1 About the Project

1.1 PMV-PPD Prediction

Package to calculate several thermal comfort indices (e.g. PMV, PPD, SET, adaptive) and convert physical variables.

Pythermalcomfort is an open-source Python library that provides a set of tools and functions for assessing thermal comfort in indoor and outdoor environments. The library implements various thermal comfort models, standards, and indices, including ASHRAE-55, ISO 7730, and PMV/PPD. The library is developed and maintained by the Center for the Built Environment (CBE) at the University of California, Berkeley.

The pythermalcomfort library includes the following functionalities:

1. Calculate Thermal Comfort Indices: The library includes functions to calculate various thermal comfort indices, such as Predicted Mean Vote (PMV), Predicted Percentage of Dissatisfied (PPD), and Thermal Sensation Vote (TSV).
2. Estimate Occupant Comfort: The library can estimate the comfort level of occupants based on their thermal sensation votes, clothing insulation, and metabolic rates.
3. Analyze Environmental Parameters: The library includes functions to analyze various environmental parameters, such as air temperature, radiant temperature, air velocity, and humidity.
4. Calculate Energy Consumption: The library can estimate energy consumption for heating and cooling systems based on thermal comfort indices and environmental parameters.
5. Compare Thermal Comfort Models: The library includes functions to compare different thermal comfort models and standards, such as ASHRAE-55 and ISO 7730, and evaluate their performance in different contexts.

1.2 Technical Aspects of this Project

This Project aims to provide a set of tools and functions for evaluating thermal comfort in indoor and outdoor environments, as well as estimating energy consumption for heating and cooling systems.

The library includes various thermal comfort models, such as ASHRAE-55, ISO 7730, and PMV/PPD, which are widely used in the building industry. These models use environmental parameters such as air temperature, radiant temperature, air velocity, and humidity, as well as human factors such as metabolic rate and clothing insulation, to predict the thermal sensation of occupants. The library can estimate the comfort level of occupants based on their thermal sensation votes, clothing insulation, and metabolic rates.

In addition to thermal comfort evaluation, the pythermalcomfort library includes tools for data visualization, such as heatmaps, scatter plots, and histograms. The library can also estimate energy consumption for heating and cooling systems based on thermal comfort indices and environmental parameters.

The library is well-documented, and the API is easy to use, making it accessible to a wide range of users, including architects, engineers, building managers, and researchers. The library is regularly updated with the latest research and standards in the field of thermal comfort, ensuring that it remains relevant and up-to-date.

1.3 Sample Output

```
Last login: Sun Apr 9 23:15:20 on ttys000
(base) sujays@Sujays-MacBook-Air ~ % cd Desktop
(base) sujays@Sujays-MacBook-Air Desktop % cd pythermalcomfort_ml
(base) sujays@Sujays-MacBook-Air pythermalcomfort_ml % ls
AUTHORS.rst           CODE_OF_CONDUCT.md      MANIFEST.in          docs           setup.py        tox.ini
CHANGELOG.rst         CONTRIBUTING.rst       README.rst          examples        src
CITATION.bib          LICENSE               ci                  setup.cfg      tests
(base) sujays@Sujays-MacBook-Air pythermalcomfort_ml % cd examples
(base) sujays@Sujays-MacBook-Air examples % ls
Untitled.ipynb         calc_adaptive_EN.py      calc_pmv_ppd.py    setup.py        tox.ini
calc_adaptive_ASHRAE.py calc_phs.py          calc_set_tmp.py   calc_utci.py   template-SI.csv
(base) sujays@Sujays-MacBook-Air examples % python3 calc_pmv_ppd.py
{'pmv': -0.55, 'ppd': 11.4}
pmv=-0.55, ppd=11.4%
{'pmv': -0.38, 'ppd': 8.1}
    tdb  tr   v  rh  met  clo   vr  clo_d  pmv  ppd
0   25  25  0.15  50  1.0   1  0.15  1.000  0.43  13.2
1   26  25  0.15  50  1.3   1  0.24  0.908  0.63  13.2
2   27  25  0.15  50  1.6   1  0.33  0.850  0.88  18.5
3   28  25  0.15  50  1.9   1  0.42  0.811  0.98  25.4
4   29  25  0.15  50  2.2   1  0.51  0.782  1.19  34.9
2.3251330852508545
0.010466978977783203
(base) sujays@Sujays-MacBook-Air examples %
```

2 Your observation on project and possible additional functionalities

2.1 Functionality 1 - Addition of more Attributes

The addition of more attributes to the calculation of Predicted Mean Vote (PMV) and Predicted Percentage of Dissatisfied (PPD) can enhance the accuracy of thermal comfort evaluations, particularly when using ISO and ASHRAE indices. However, the practical implementation of such enhancements may be constrained by factors such as the availability and cost of sensors and other equipment required to measure and record the additional attributes. Therefore, it is important to carefully consider the benefits and costs of adding more attributes and ensure that the proposed enhancements are feasible and cost-effective for the specific context and requirements of the project.

2.2 Functionality 2 - Machine Learning Aspect

The current version of the repository calculates Predicted Mean Vote (PMV) and Predicted Percentage of Dissatisfied (PPD) and provides this information to the client. However, there is potential to enhance the functionality of the repository by incorporating additional factors that can influence thermal comfort, such as the operation of fans, air conditioning, and windows and doors for airflow, as well as clothing insulation predictions for different types of clothing, and the presence of individuals with medical conditions or of varying ages. By incorporating these factors, the repository can provide a more comprehensive and personalized assessment of thermal comfort that better reflects the needs and preferences of the occupants. However, the incorporation of additional factors may also require additional sensors and data collection equipment, as well as more complex algorithms for analysis and prediction, which may increase the complexity and cost of the system. Therefore, it is important to carefully evaluate the costs and benefits of incorporating additional factors and ensure that any proposed enhancements are feasible, reliable, and effective in meeting the needs of the clients and end-users.

2.3 Functionality 3 - Semi-Supervised Learning

The Usage of Semi-Supervised Learning Comes after the Prediction of PMV and PPD once the Prediction is done we use the Semi-Supervised Learning to Classify Is the Thermal sensation is good-enough or not or if its an yes/no then on which level and by this we can conclude this DataStream Application from the Sensors give us an enumerate approach to predict and classify for the persons living inside the room.

The application of semi-supervised learning can be a valuable addition to the existing prediction of Predicted Mean Vote (PMV) and Predicted Percentage of Dissatisfied (PPD) in thermal comfort assessment. After the prediction of PMV and PPD, semi-supervised learning can be used to classify the thermal sensation as good enough or not and, if not, at which level. This can provide a more detailed and nuanced evaluation of thermal comfort that is tailored to the specific needs and preferences of the occupants. The incorporation of semi-supervised learning can also allow for a more dynamic and responsive system that can adapt to changes in the environment and the occupants over time.

By leveraging data from sensors in real-time, the data stream application can provide a more accurate and up-to-date assessment of thermal comfort, which is particularly valuable in dynamic and unpredictable environments. However, the implementation of semi-supervised learning may require additional computational resources and expertise in machine learning, which can increase the complexity and cost of the system. Therefore, it is important to carefully consider the costs and benefits of incorporating semi-supervised learning and ensure that the proposed enhancements are feasible, reliable, and effective in meeting the needs of the clients and end-users.

3 My Feature

```
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
import numpy as np
import pandas as pd
import time
from google.colab import drive
from keras.layers import LSTM
from keras.models import Sequential
from keras.layers.core import Dense, Activation, Dropout
from sklearn.preprocessing import MinMaxScaler
import keras
from sklearn import datasets
from sklearn import metrics
from sklearn.neural_network import MLPClassifier
from sklearn.neural_network import MLPRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
22s [7] drive.mount('/content/gdrive')
Mounted at /content/gdrive
[17] df = pd.read_csv("db_measurements_v2.1.0.csv")
<ipython-input-17-41cebf8d01>:1: DtypeWarning: Columns (5,35,36) have mixed types. Specify dtype option on import or set low_memory=False.
df = pd.read_csv("db_measurements_v2.1.0.csv")
[18] df.columns
0s Index(['index', 'record_id', 'building_id', 'timestamp', 'season',
       'subject_id', 'age', 'gender', 'ht', 'wt', 'ta', 'ta_h', 'ta_m', 'ta_l',
       'tr', 'tr_h', 'tr_m', 'tr_l', 'rh', 'rh_h', 'rh_m', 'rh_l',
       'vel', 'vel_h', 'vel_m', 'vel_l', 'met', 'clo', 'activity_10', 'activity_20',
       'activity_30', 'activity_60', 'thermal_sensation',
       'thermal_acceptability', 'thermal_preference', 'thermal_comfort',
       'air_movement_acceptability', 'air_movement_preference',
       'blind_curtain', 'fan', 'window', 'door', 'heater', 't_out', 'rh_out',
       't_mot_isd'],
      dtype='object')
[33] considered = df[['ta', 'tr', 'vel', 'rh', 'met', 'clo']]
predictors = df[['pmv', 'ppd', 'pmv_ce', 'ppd_ce']]
classifiers = df[['thermal_acceptability', 'thermal_preference', 'thermal_comfort']]
```

```
[33] considered = df[['ta', 'tr', 'vel', 'rh', 'met', 'clo']]
predictors = df[['pmv', 'ppd', 'pmv_ce', 'ppd_ce']]
classifiers = df[['thermal_acceptability', 'thermal_preference', 'thermal_comfort']]
0s
[34] df[['thermal_acceptability', 'thermal_preference', 'thermal_comfort']]
   thermal_acceptability thermal_preference thermal_comfort
0   acceptable           cooler            4.0
1 unacceptable          cooler            2.0
2   acceptable          no change           5.0
3   acceptable          no change           5.0
4   acceptable          no change           5.0
...
109028      NaN        NaN        NaN
109029      NaN        NaN        NaN
109030      NaN        NaN        NaN
109031      NaN        NaN        NaN
109032      NaN        NaN        NaN
109033 rows × 3 columns
0s
[35] df[['pmv', 'ppd', 'pmv_ce', 'ppd_ce']]
   pmv  ppd  pmv_ce  ppd_ce
0  0.50  10.2   0.38    8.0
1  0.40   8.4   0.40    8.4
2 -0.07   5.1  -0.07    5.1
3  0.31   7.0   0.14    5.4
4  0.05   5.0  -0.06    5.1
...
109028      NaN      NaN      NaN      NaN
```

The screenshot shows a Google Colab notebook titled "PMV_PPD.ipynb". The code cell [50] displays the data types of the columns in the DataFrame:

```
[50] new_df.dtypes
```

The output shows:

Column	Dtype
ta	float64
tr	float64
vel	float64
rh	float64
met	float64
clo	float64
pmv	float64
ppd	float64
ppd_ce	float64
thermal_acceptability	object
thermal_preference	object
thermal_comfort	float64
dtype	object

Code cell [51] shows the value counts for the "thermal_acceptability" column:

```
[51] new_df["thermal_acceptability"].value_counts()
```

The output shows:

Category	Count
acceptable	5582
unacceptable	800

Code cell [54] shows the value counts for the "thermal_preference" column:

```
[54] new_df["thermal_preference"].value_counts()
```

The output shows:

Category	Count
no change	3698
cooler	1581
warmer	1103

Code cell [56] shows the value counts for the "thermal_comfort" column:

```
[56] new_df["thermal_comfort"].value_counts()
```

The output shows:

Category	Count
5.0	3193
4.0	1245
6.0	989
3.0	757
2.0	205
1.0	73

The screenshot continues the execution of the notebook "PMV_PPD.ipynb". The code cell [56] shows the value counts for the "thermal_comfort" column:

```
[56] new_df["thermal_comfort"].value_counts()
```

The output shows:

Category	Count
5.0	3193
4.0	1245
6.0	989
3.0	757
2.0	205
1.0	73

Code cell [61] initializes variables X and y:

```
[61] X = new_df[["ta","tr","vel","rh","met","clo"]]
y = new_df[["pmv"]]
```

Code cell [63] imports train_test_split and splits the data:

```
[63] from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,random_state=1)
```

Code cell [69] imports SVR and fits the model:

```
[69] from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error,r2_score
regressor = SVR(kernel = 'rbf')
regressor.fit(X_train, y_train)
y_pred = regressor.predict(X_test)
```

Code cell [71] saves the model to disk:

```
[71] # save the model to disk
import pickle
filename = 'pmv_model.sav'
pickle.dump(regressor, open(filename, 'wb'))
```

The notebook indicates completion at 11:11 PM.

```

[72] X1 = new_df[['ta','tr','vel','rh','met','clo']]
y1 = new_df[['ppd']]

[73] from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
sc_Y = StandardScaler()
X1 = sc_X.fit_transform(X1)
y1 = sc_Y.fit_transform(y1)

[75] from sklearn.model_selection import train_test_split
X_train1, X_test1, y_train1, y_test1 = train_test_split(X1, y1, random_state=1)

[86] from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score
regressor = SVR(kernel = 'rbf')
regressor.fit(X_train1, y_train1)
y_pred1 = regressor.predict(X_test1)
mean_squared_error(y_pred1,y_test1)
r2_score(y_pred1,y_test1)

[87] # save the model to disk
import pickle
filename = 'ppd_model.sav'
pickle.dump(regressor, open(filename, 'wb'))

loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.predict(X_test1)
print(r2_score(result,y_test1))

0.9746474966198946

```

UNSUPERVISED LEARNING FOR THESE ATTRIBUTES

```

[78] X2 = new_df[['ta','tr','vel','rh','met','clo','pmv','ppd']]

```

```

[78] X2 = new_df[['ta','tr','vel','rh','met','clo','pmv','ppd']]

[82] from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled_df_kmeans = scaler.fit_transform(X2)

[83] from sklearn.cluster import KMeans
kmeans_model = KMeans(
    n_clusters=3, init='random',
    n_init=10, max_iter=300,
    tol=1e-04, random_state=0
)
y_km = kmeans_model.fit_predict(X2)
clusters = kmeans_model.labels_
clusters

[88] # save the model to disk
import pickle
filename = 'unsup_kMeans_model.sav'
pickle.dump(kmeans_model, open(filename, 'wb'))

[92] pd.DataFrame(clusters).value_counts()

0    3972
1    1379
2    1131
dtype: int64

[94] X2

```

	ta	tr	vel	rh	met	clo	pmv	ppd
0	22.3	22.950000	0.03	61.0	1.706485	0.95	0.50	10.2
1	23.0	24.104286	0.08	59.0	1.109215	1.07	0.40	8.4
2	22.0	22.091429	0.04	61.0	1.211604	0.88	-0.07	5.1

A screenshot of a Google Colab notebook titled "PMV_PPD.ipynb". The code cell at line 246 shows the creation of three clusters (cluster0, cluster1, cluster2) from a DataFrame. The clusters are defined as dictionaries mapping column names to lists of values. The code then creates DataFrames for each cluster. The output of the cell shows the first few rows of the cluster0 DataFrame.

```
[246]: cluster0 = []
cluster1 = []
cluster2 = []
cluster0 = {
    'ta': [0],
    'tr': [0],
    'vel': [0],
    'rh': [0],
    'met': [0],
    'clo': [0],
    'pmv': [0],
    'ppd': [0],
    'pmv_ce': [0],
    'ppd_ce': [0],
    'thermal_acceptability': [0],
    'thermal_preference': [0],
    'thermal_comfort': [0],
}
cluster0 = pd.DataFrame(cluster0)

cluster1 = [
    'ta': [0],
    'tr': [0],
    'vel': [0],
    'rh': [0],
    'met': [0],
    'clo': [0],
]
cluster1 = pd.DataFrame(cluster1)

cluster2 = [
    'ta': [0],
    'tr': [0],
    'vel': [0],
    'rh': [0],
    'met': [0],
    'clo': [0],
    'pmv': [0],
    'ppd': [0],
    'pmv_ce': [0],
    'ppd_ce': [0],
    'thermal_acceptability': [0],
    'thermal_preference': [0],
    'thermal_comfort': [0],
]
cluster2 = pd.DataFrame(cluster2)
```

The output of the cell is:

```
[95]: new_df.iloc[0]
```

	ta	tr	vel	rh	met	clo	pmv	ppd	pmv_ce	ppd_ce	thermal_acceptability	thermal_preference	thermal_comfort
Name: 0, dtype: object	22.3	22.95	0.03	61.0	1.765625	0.95	0.5	10.2	0.38	8.0	1	cooler	4.0

A screenshot of a Google Colab notebook titled "PMV_PPD.ipynb". The code cell at line 246 continues the creation of three clusters (cluster0, cluster1, cluster2) from a DataFrame. The clusters are defined as dictionaries mapping column names to lists of values. The code then creates DataFrames for each cluster. The output of the cell shows the first few rows of the cluster0 DataFrame.

```
[246]: cluster0 = []
cluster1 = []
cluster2 = []
cluster0 = {
    'ta': [0],
    'tr': [0],
    'vel': [0],
    'rh': [0],
    'met': [0],
    'clo': [0],
    'pmv': [0],
    'ppd': [0],
    'pmv_ce': [0],
    'ppd_ce': [0],
    'thermal_acceptability': [0],
    'thermal_preference': [0],
    'thermal_comfort': [0],
}
cluster0 = pd.DataFrame(cluster0)

cluster1 = [
    'ta': [0],
    'tr': [0],
    'vel': [0],
    'rh': [0],
    'met': [0],
    'clo': [0],
]
cluster1 = pd.DataFrame(cluster1)

cluster2 = [
    'ta': [0],
    'tr': [0],
    'vel': [0],
    'rh': [0],
    'met': [0],
    'clo': [0],
    'pmv': [0],
    'ppd': [0],
    'pmv_ce': [0],
    'ppd_ce': [0],
    'thermal_acceptability': [0],
    'thermal_preference': [0],
    'thermal_comfort': [0],
]
cluster2 = pd.DataFrame(cluster2)
```

The output of the cell is:

```
[95]: new_df.iloc[0]
```

	ta	tr	vel	rh	met	clo	pmv	ppd	pmv_ce	ppd_ce	thermal_acceptability	thermal_preference	thermal_comfort
Name: 0, dtype: object	22.3	22.95	0.03	61.0	1.765625	0.95	0.5	10.2	0.38	8.0	1	cooler	4.0

At the bottom of the cell, there is a warning message:

```
<ipython-input-103-ddbf279392c4>: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use par
cluster0 = cluster0.append(new_df.iloc[[1]])
cluster0 = cluster0.append(new_df.iloc[[1]])
<ipython-input-103-ddbf279392c4>: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use par
cluster0 = cluster0.append(new_df.iloc[[1]])
<ipython-input-103-ddbf279392c4>: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use par
cluster0 = cluster0.append(new_df.iloc[[1]])
```

The screenshot shows a Jupyter Notebook titled "PMV_PPD.ipynb" running on Google Colab. The notebook contains a DataFrame named "df" with 109032 rows and 22 columns. The columns are: index, record_id, building_id, timestamp, season, subject_id, age, gender, ht, wt, t_out, rh_out, t_out_isd, rh_out_isd, set, pmv, ppd, and p. Below the DataFrame, there are two code cells:

```
0 [187] second = pd.read_csv("two.csv")
0 [189] second.dtypes
```

The first cell reads the CSV file "two.csv" into a variable "second". The second cell prints the data types for each column in "second". The output shows that most columns have data type "float64", while "subject_id", "age", "gender", "ht", "wt", "t_out", "rh_out", "t_out_isd", "rh_out_isd", "set", "pmv", "ppd", and "p" have data type "int64".

```

File Edit View Insert Runtime Tools Help All changes saved
Files + Code + Text
[107] second = pd.read_csv("two.csv")
[109] second.dtypes
[110] second["thermal_comfort"].value_counts()
[112] second.corr()

```

Detailed description: This screenshot shows a Jupyter Notebook interface on Google Colab. The code cell at the top reads a CSV file named 'two.csv' into a DataFrame called 'second'. It then displays the data types for each column. The next cell shows the value counts for the 'thermal_comfort' column, which has 672 entries ranging from 0.0 to 5.0. The final cell calculates the correlation matrix for all columns in 'second'. The interface includes a sidebar with file navigation, a code editor, and a text editor.

```

File Edit View Insert Runtime Tools Help All changes saved
Files + Code + Text
[123] second_df = second.drop(0,axis=0)
[124] second_df = second.drop("Unnamed: 0",axis=1)
[125] X = second_df.iloc[:,1:10]
[126] y = second_df["thermal_acceptability"]
[127] from sklearn.linear_model import LogisticRegression
[128] from sklearn.model_selection import train_test_split
[129] from sklearn.metrics import accuracy_score
[130] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
[131] # create a logistic regression model
[132] lr = LogisticRegression()
[133] # fit the model with the training data
[134] lr.fit(X_train, y_train)
[135] # make predictions on the testing data
[136] y_pred = lr.predict(X_test)
[137] # evaluate the model's accuracy
[138] accuracy = accuracy_score(y_test, y_pred)
[139] print("Accuracy:", accuracy)

```

Detailed description: This screenshot shows a Jupyter Notebook interface on Google Colab. The code cell at the top drops the first row and the 'Unnamed: 0' column from the DataFrame 'second'. The next cell splits the data into training and testing sets. Subsequent cells import the 'LogisticRegression' class from scikit-learn, create an instance of it, fit it to the training data, and then predict on the testing data. The accuracy of the model is printed. A warning message about convergence is visible. The interface includes a sidebar with file navigation, a code editor, and a text editor.

A screenshot of a Google Colab notebook titled "PMV_PPD.ipynb". The code cell contains Python code for training a Random Forest classifier on a dataset. The output shows the accuracy of the model as 0.8333333333333334. The code then saves the trained model to a file named "thermal_acc_model.sav".

```
[132] from sklearn.ensemble import RandomForestClassifier
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import accuracy_score

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

      # create a Random Forest classifier with 100 trees
      rf = RandomForestClassifier(n_estimators=100)

      # fit the model with the training data
      rf.fit(X_train, y_train)

      # make predictions on the testing data
      y_pred = rf.predict(X_test)

      # evaluate the model's accuracy
      accuracy = accuracy_score(y_test, y_pred)
      print("Accuracy:", accuracy)

      Accuracy: 0.8333333333333334

      # save the model to disk
      import pickle
      filename = 'thermal_acc_model.sav'
      pickle.dump(rf, open(filename, 'wb'))
```

A screenshot of a Google Colab notebook titled "PMV_PPD.ipynb". The code cell contains Python code for training a Logistic Regression model on the "thermal_comfort" data. The output shows the accuracy of the model as 0.5. A warning message indicates that the optimization algorithm (lbfgs) failed to converge, and it suggests increasing the number of iterations or scaling the data. The code then saves the trained model to a file named "thermal_comfort_model.sav".

```
[134] X1 = second_df.iloc[:,10]
      y1 = second_df["thermal_comfort"]

      from sklearn.linear_model import LogisticRegression
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import accuracy_score

      X_train, X_test, y_train, y_test = train_test_split(X1, y1, test_size=0.3, random_state=42)

      # create a logistic regression model
      lr1 = LogisticRegression()

      # fit the model with the training data
      lr1.fit(X_train, y_train)

      # make predictions on the testing data
      y_pred = lr1.predict(X_test)

      # evaluate the model's accuracy
      accuracy = accuracy_score(y_test, y_pred)
      print("Accuracy:", accuracy)

      Accuracy: 0.5
      /usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
      STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

      Increase the number of iterations (max_iter) or scale the data as shown in:
      https://scikit-learn.org/stable/modules/preprocessing.html
      Please also refer to the documentation for alternative solver options:
      https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
      n_iter_1 = _check_optimize_result()

      # save the model to disk
      import pickle
      filename = 'thermal_comfort_model.sav'
      pickle.dump(lr1, open(filename, 'wb'))
```

4 References

https://github.com/sujaykumarmag/pythermalcomfort_ml

<https://www.softxjournal.com/action/showPdf?pii=S2352-7110%2820%2930245-4>

<https://cbe-berkeley.gitbook.io/thermal-comfort-tool/>

<https://pythermalcomfort.readthedocs.io/en/latest/usage.html>